

Biomedical Data Analysis - Lab 4

Group members:

- Enrico Maria Marinelli
- Francesco Pio Capoccello
- Juras Lukaševičius

Libraries used in practical session:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score, roc_curve, auc
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, classification_report
from sklearn.metrics import confusion_matrix
from sklearn.impute import KNNImputer
from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import cross_val_score
from IPython.display import HTML
from scipy.stats import pearsonr
import regex as re
```

Block I - Data loading and basic description

We load the dataset using the pandas' load_csv function, remember to specify semicolon (;) as the separator

```
df=pd.read_csv("../BiomedicalDataScience_LAB04//Data/
inadvance_synth.csv",sep=";")
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 38416 entries, 0 to 38415
```

```
Data columns (total 22 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	38416 non-null	int64
1	age	38416 non-null	int64
2	barthel	10218 non-null	float64
3	charlson	38339 non-null	float64
4	codidiagingreso	27038 non-null	object
5	codservicioreal	38416 non-null	object
6	creatinina	30465 non-null	float64

7	drg	34776	non-null	float64
8	estancias	38416	non-null	float64
9	glucosa	29188	non-null	float64
10	hematocrito	30114	non-null	float64
11	leucocitos	30114	non-null	float64
12	metastatic_tumor	38339	non-null	float64
13	num_grupoact3_HOSP	17415	non-null	float64
14	numurgenciasprevias	37864	non-null	float64
15	potasio	29770	non-null	float64
16	proteina_c_reactiva	20241	non-null	float64
17	rdw_cv	30082	non-null	float64
18	rdw_sd	30082	non-null	float64
19	sodio	30274	non-null	float64
20	urea	24156	non-null	float64
21	label	38416	non-null	float64

dtypes: float64(18), int64(2), object(2)
memory usage: 6.4+ MB

1. What is the size of the dataframe?

```
# Drop the "Unnamed" column
df = df.iloc[:,1:]

df.shape

(38416, 21)
```

2. What is the mean age?

```
mean = df.age.mean()
print(mean)

79.4033215326947
```

3. What is the age standard deviation (std)?

```
std=df.age.std()
print(std)

8.361737254115944
```

4. Which is the variable with the most amount of missing values? Can you list the name of the variables, sorting them by number of missing values?

```
def nullvalues(df=df):
    nullvalues = df.isnull().sum()
    features_sorted = nullvalues.sort_values(ascending=False)
```

```

    return features_sorted
print(nullvalues())

barthel      28198
num_grupoact3_HOSP  21001
proteina_c_reactiva  18175
urea         14260
codidiagingreso  11378
glucosa      9228
potasio      8646
rdw_sd       8334
rdw_cv       8334
hematocrito  8302
leucocitos   8302
sodio        8142
creatinina   7951
drg          3640
numurgenciasprevias  552
metastatic_tumor  77
charlson     77
age          0
estancias    0
codservicioreal  0
label        0
dtype: int64

```

5. Name which are the categorical variables.

```

categorical_variables = df.select_dtypes(include=['object'])
categorical_variables_names = categorical_variables.columns.tolist()
print(categorical_variables_names)

['codidiagingreso', 'codservicioreal']

def find_cat(x):
    if isinstance(x, str):
        pattern = r'[0-9]+.[0-9]*'
        test = not (re.match(pattern, x))
        # print(f"{x}: {test}") # Add this line for debugging
        return test
    # print(f"{x}: False") # Add this line for debugging
    return False

def
categorical_enumeration(cat, categorical_variables=categorical_variables,
    printing=True):

    if printing:
        print('Evaluation of categorical variable "' + cat + "'")

```

```

categories =
categorical_variables[categorical_variables[cat].apply(find_cat)]
if cat == categorical_variables_names[0]:
    categories =
categorical_variables[categorical_variables[cat].apply(find_cat)]
else:
    categories = categorical_variables[cat].astype(str)

unique, count = np.unique(categories, return_counts=True)

if printing:
    print('Number of categories for "' + cat + '":', len(unique))

combined = np.column_stack((unique, count))
sorted_combined = combined[combined[:, 1].argsort()[::-1]]

if printing:
    for i in sorted_combined:
        if isinstance(i[0], float):
            print("Class %f has %d samples" % (i[0], i[1]))
        else:
            print("Class " + str(i[0]) + " has %d samples" % i[1])

    return np.array(sorted_combined)

categories_codidiagingreso =
categorical_enumeration(categorical_variables_names[0])

```

Evaluation of categorical variable "codidiagingreso"

Number of categories for "codidiagingreso": 299

Class HMUR has 507 samples
 Class HNEM has 382 samples
 Class HCAR has 375 samples
 Class V14.8 has 245 samples
 Class HMIN has 215 samples
 Class HONC has 214 samples
 Class V15.82 has 195 samples
 Class HNER has 183 samples
 Class V14.0 has 181 samples
 Class V58.61 has 121 samples
 Class V14.1 has 105 samples
 Class HHEM has 78 samples
 Class M81403 has 76 samples
 Class V58.67 has 71 samples
 Class HURO has 65 samples
 Class V14.6 has 63 samples
 Class HMDG has 54 samples
 Class V45.01 has 42 samples

```
Class M96923 has 1 samples
Class M98233 has 1 samples
Class M97313 has 1 samples
Class M98003 has 1 samples
Class M98013 has 1 samples
Class M84416 has 1 samples
Class M83233 has 1 samples
Class V18.8 has 1 samples
Class V10.82 has 1 samples
Class V17.1 has 1 samples
Class M98103 has 1 samples
Class HSEP has 1 samples
Class HUEG has 1 samples
Class HUML has 1 samples
Class V13.29 has 1 samples
Class V13.09 has 1 samples
Class V12.50 has 1 samples
Class M80203 has 1 samples
Class M80206 has 1 samples
Class V12.41 has 1 samples
Class V12.02 has 1 samples
Class V10.44 has 1 samples
Class V10.04 has 1 samples
Class V10.43 has 1 samples
Class M80906 has 1 samples
Class V10.41 has 1 samples
Class M81303 has 1 samples
Class V10.09 has 1 samples
Class V10.07 has 1 samples
Class M81703 has 1 samples
Class M82103 has 1 samples
Class M82403 has 1 samples
Class M82606 has 1 samples
Class M82613 has 1 samples
Class M82630 has 1 samples
Class E812.1 has 1 samples
```

```
categories_codservicioreal =
categorical_enumeration(categorical_variables_names[1])
```

Evaluation of categorical variable "codservicioreal"

Number of categories for "codservicioreal": 53

```
Class HMUR has 8233 samples
Class HNEM has 5422 samples
Class HCAR has 5124 samples
Class HMIN has 3633 samples
Class HNER has 2341 samples
Class HMDG has 2222 samples
Class HONC has 2041 samples
Class HUHP has 1159 samples
```

```
Class HURO has 1112 samples
Class HCDG has 1007 samples
Class HCOT has 741 samples
Class HHEM has 546 samples
Class HNEF has 522 samples
Class HCVA has 502 samples
Class HMDH has 472 samples
Class HCEP has 420 samples
Class HNCG has 381 samples
Class HUMI has 351 samples
Class HUEI has 332 samples
Class HCCV has 308 samples
Class HREU has 252 samples
Class HORL has 179 samples
Class HREA has 157 samples
Class HCTO has 154 samples
Class HCIR has 146 samples
Class HECR has 126 samples
Class HCLP has 71 samples
Class HSEP has 60 samples
Class HGIN has 55 samples
Class HUTP has 50 samples
Class HCMX has 49 samples
Class HUML has 44 samples
Class HCPL has 43 samples
Class HUMM has 31 samples
Class HUEM has 28 samples
Class HOFT has 22 samples
Class HRHB has 16 samples
Class HTRE has 15 samples
Class HULM has 13 samples
Class HRXD has 8 samples
Class HQUE has 5 samples
Class HALE has 4 samples
Class HDER has 4 samples
Class HUEG has 4 samples
Class HURQ has 2 samples
Class HTRA has 2 samples
Class HREM has 1 samples
Class HRER has 1 samples
Class HRQU has 1 samples
Class HRTE has 1 samples
Class HUDO has 1 samples
Class HUHT has 1 samples
Class HONG has 1 samples
```

In addition to the existing `codidiagingreso` and `codservicioreal` categorical variables including codes in the format

str

is `metastatic_tumor`, since its only values are 0 and 1.

6. Extract the 'label' column to another variable. How many positive cases there are? And negatives?

```
label_var = df['label']
positive = df[label_var==1].shape[0]
negative = df[label_var==0].shape[0]
print(f"Positive cases: {positive}")
print(f"Negative cases: {negative}")
```

```
Positive cases: 13431
Negative cases: 24985
```

Block II. Data preprocessing

1. Train-Test Split

How many samples have each set after the split?

Split the dataset in two: train (80%) and test (20%). Use a seed to allow replication.

```
seed = 42

X = df.drop('label', axis=1)
y = df['label']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=seed)

rows_train = X_train.shape[0]
rows_test = X_test.shape[0]
cols = X_train.shape[1]

print(f"Train samples: {rows_train}")
print(f"Test samples: {rows_test}")
print(f"Features: {cols}")

Train samples: 30732
Test samples: 7684
Features: 20
```

2. Categorical Variables

Implement the method to deal with categorical variables. Briefly explain the chosen alternative. What would happen if a variable in the test set contains a category that doesn't exist on the train set? How would you deal with this situation?

Categories: `codidiagingreso` and `codservicioreal`

Is required to have numerical values only in order to let a *Machine Learning Model* perform the training phase and, subsequently, the inference procedure.

It is possible to deal with categorical variables through a conversion to scalar numerical variables or create a vector *One hot encoded* substituting each `string` variable, that will be *unwinded* into the dataframe creating a new feature for each of them.

The first approach is followed in the previous defined function `categorical_enumeration` and the following `categorical_evaluation`, accordingly to the name, it is based on frequency (enumeration) of the appearing variables.

Furthermore, this approach will be abandoned in favour of the second one.

```
def categorical_evaluation(df,alpha=10):

    #
    # alpha variable represents how much weight the feature
    "metastatic tumor"
    # has on the categorical variable.
    # This choice is given by the fact that features might need to be
    more spread out.
    #
    # NOTE: alpha is not an tunable hyperparameter, but might need to
    be suggested by a
    # medic or someone who has knowledge in the field.
    # It has been set to 10 by default to spread more the low
    frequency variables,
    # but might be not a correct assumption

    dataset = df.copy()
    cv = dataset.select_dtypes(include=['object'])
    for cat in cv.columns.tolist():
        enum =
        categorical_enumeration(cat=cat,categorical_variables=cv,printing=False)

        # The ranking process starts here
        mapping = {}
        for i,e in enumerate(enum[:,0]):

            # Increase the feature "importance rank" by a slight
            correlation with metastatic tumor variable
            # only for low frequency variables
            if enum[i,1]<5:
                increase =
                np.sum(np.logical_and(dataset['metastatic_tumor']==1,dataset[cat]==e))

                # This increase is regulate by the aforementioned
```



```

alpha parameter
    enum[i,1] += enum[i,1]*(alpha*increase)

    enum[i,1] /= np.sum(enum[:,1]) # Represent the percentage
of appearance in the given dataset (instead of the frequency)
    mapping[e] = enum[i,1]
    dataset[cat] = dataset[cat].map(mapping)
    dataset[cat] = dataset[cat].astype(float,errors='ignore') #
'coerce' will handle non-convertible values
    # filter =
np.logical_and(dataset.isnull(),np.logical_not(df.isnull()))
    # print(df[filter])

```

```

#nullvalues(dataset)
return dataset

```

```

X_train_categorical_mapped = categorical_evaluation(X_train)
X_test_categorical_mapped = categorical_evaluation(X_test)

```

```

X_train_categorical_mapped.info()

```

```

<class 'pandas.core.frame.DataFrame'>

```

```

Index: 30732 entries, 29147 to 15795

```

```

Data columns (total 20 columns):

```

#	Column	Non-Null Count	Dtype
0	age	30732 non-null	int64
1	barthel	8222 non-null	float64
2	charlson	30673 non-null	float64
3	codidiagingreso	1905 non-null	float64
4	codservicioreal	30732 non-null	float64
5	creatinina	24410 non-null	float64
6	drg	27841 non-null	float64
7	estancias	30732 non-null	float64
8	glucosa	23372 non-null	float64
9	hematocrito	24144 non-null	float64
10	leucocitos	24144 non-null	float64
11	metastatic_tumor	30673 non-null	float64
12	num_grupoact3_HOSP	13964 non-null	float64
13	numurgenciasprevias	30284 non-null	float64
14	potasio	23860 non-null	float64
15	proteina_c_reactiva	16210 non-null	float64
16	rdw_cv	24120 non-null	float64
17	rdw_sd	24120 non-null	float64
18	sodio	24269 non-null	float64
19	urea	19388 non-null	float64

```
dtypes: float64(19), int64(1)
memory usage: 4.9 MB
```

3.1 Dealing with Categorical Variables missing

Here we have decided to discard the categorical mapping approach in favour of an **encoding approach**, that let us achieve better results

and is more robust with respect to data acquired.

```
median_train = X_train['codidiagngreso'].apply(pd.to_numeric,
errors='coerce').median()
```

```
X_train['codidiagngreso'].fillna(median_train, inplace=True)
```

```
X_test['codidiagngreso'].fillna(median_train, inplace=True)
```

```
X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 30732 entries, 29147 to 15795
```

```
Data columns (total 20 columns):
```

#	Column	Non-Null Count	Dtype
0	age	30732 non-null	int64
1	barthel	8222 non-null	float64
2	charlson	30673 non-null	float64
3	codidiagngreso	30732 non-null	object
4	codservicioreal	30732 non-null	object
5	creatinina	24410 non-null	float64
6	drg	27841 non-null	float64
7	estancias	30732 non-null	float64
8	glucosa	23372 non-null	float64
9	hematocrito	24144 non-null	float64
10	leucocitos	24144 non-null	float64
11	metastatic_tumor	30673 non-null	float64
12	num_grupoact3_HOSP	13964 non-null	float64
13	numurgenciasprevias	30284 non-null	float64
14	potasio	23860 non-null	float64
15	proteina_c_reactiva	16210 non-null	float64
16	rdw_cv	24120 non-null	float64
17	rdw_sd	24120 non-null	float64
18	sodio	24269 non-null	float64
19	urea	19388 non-null	float64

```
dtypes: float64(17), int64(1), object(2)
```

```
memory usage: 4.9+ MB
```

```
print(f"X_train codidiagngreso null values:
{X_train['codidiagngreso'].isnull().sum()} \n\
```

```
X_test codidiagingreso null values:  
{X_test['codidiagingreso'].isnull().sum()}")
```

```
X_train codidiagingreso null values: 0  
X_test codidiagingreso null values: 0
```

We have used the function `pd.to_numeric` to convert the values in the `codidiagingreso` column to numbers. The `errors='coerce'` option causes non-numeric values to be replaced with NaN (Not a Number) during the conversion. At the end of this operation, we get a Series of numeric values and NaN.

After converting the values to numbers (and replacing the non-numeric values with NaN), we calculate the median of the Train Set.

The median is calculated by ignoring NaN values, so you get the median only from the valid numeric values in the `codidiagingreso` column.

Then we applied this median obtained in the Training Set on the Test Set, as it was.

Category: `metastatic_tumor`

```
k_values = list(range(1, 11))  
  
mean_scores = []  
for k in k_values:  
    imputer = KNNImputer(n_neighbors=k)  
    imputed_data =  
    imputer.fit_transform(X_train[['metastatic_tumor']])  
  
    knn_regressor = KNeighborsRegressor(n_neighbors=k)  
    scores = cross_val_score(knn_regressor, imputed_data, y_train,  
cv=10, scoring='neg_mean_squared_error')  
    mean_scores.append(np.mean(scores))  
  
best_k = k_values[np.argmin(mean_scores)]  
print("The best number of neighbors is", best_k)  
The best number of neighbors is 1
```

The code performs a search for the best value of `k` for a KNN-based regression model.

It uses `cross_val_score` to evaluate the performance of the model on training data.

The process helps select the best `k` value for the KNN model based on the data provided.

```
imputer = KNNImputer(n_neighbors=1)  
metastatic_tumor_train = X_train[['metastatic_tumor']]  
  
X_train_imputed = imputer.fit_transform(metastatic_tumor_train)
```

```
X_train['metastatic_tumor'] = X_train_imputed
X_test['metastatic_tumor'].astype('object')
X_test['metastatic_tumor'].fillna('Unknown', inplace=True)
```

In this part it has been used `KNNImputer().fit_transform` to impute the data in the Training Set.

Whereas in the Test Set the data has been filled with `Unknown` label, that does not appear in the Training Set, a more detailed discussion will be held in the *sub-section* number 4.

3.2 Dealing with Numerical Values missing

Implement the method to deal with the missing values. Briefly explain the chosen alternative. What would happen if a variable without missing on the train set appears to have been missing in the test set? How would you deal with them?

```
# Test to verify consistency of flatten and reshape functions
# heavily used in the SVT algorithm later defined

# mat = X_full.shape
# vect = X_full.shape[0]*X_full.shape[1]
# mask = (X_full.copy().flatten().reshape(mat) == X_full).flatten()
# (X_full.flatten()[mask]-
X_full.copy().flatten().reshape(mat).flatten()[mask]).sum()

(X_train['age'] == 65).sum()

748

avoid_categoricals = X_test.columns.to_list()
categorical_variables_names.append('metastatic_tumor')
for i in categorical_variables_names:
    avoid_categoricals.remove(i)
avoid_categoricals

['age',
 'barthel',
 'charlson',
 'creatinina',
 'drg',
 'estancias',
 'glucosa',
 'hematocrito',
 'leucocitos',
 'num_grupoact3_HOSP',
 'numurgenciasprevias',
 'potasio',
 'proteina_c_reactiva',
 'rdw_cv',
```

```
'rdw_sd',
'sodio',
'urea']
```

SVT is a technique used in signal processing and matrix completion problems. It's primarily employed in scenarios where you have a matrix with missing or corrupted data, and you want to estimate the missing values based on the available information.

First of all we divide into two sets *Test* and *Train* the dataset in input.

Let us denote by Ω_{train} the set of observations in the train set, i.e. the values appearing in the original dataset corresponding to the train set,

similarly Ω_{test} is the set of observation in the test set. Let us denote by r_{ij} the corresponding rating.

Create a full matrix $X \in R^{n \times p}$, such that:

$$X_{i,j} = \begin{cases} r_{ij} & \text{if } (i,j) \in \Omega \\ 0 & \text{otherwise} \end{cases}$$

Then the algorithm is based on a iterative procedure for `n_max_iter`,

A *Singular Value Decomposition*, using `np.linalg.svd`,

a thresholding on the singular values based on `threshold`,

a `tolerance` that at each iteration is compared to the norm of the difference of the old matrix and the reconstructed new one,

the error is evaluated through:

the RMSE (root mean square error)

$$\text{RMSE} = \sqrt{\frac{1}{\text{card}(\Omega_{\text{test}})} \sum_{(i,j) \in \Omega_{\text{test}}} (r_{ij} - r_{ij}^{\text{pred}})^2}$$

and the Pearson correlation coefficient ρ , using `scipy.stats.pearsonr`:

$$\rho = \frac{\sum_{(i,j) \in \Omega_{\text{test}}} (r_{ij} - \bar{r})(r_{ij}^{\text{pred}} - \bar{r}^{\text{pred}})}{\sqrt{\sum_{(i,j) \in \Omega_{\text{test}}} (r_{ij} - \bar{r})^2} \sqrt{\sum_{(i,j) \in \Omega_{\text{test}}} (r_{ij}^{\text{pred}} - \bar{r}^{\text{pred}})^2}}$$

where

$$\bar{r} = \frac{1}{\text{card}(\Omega_{\text{test}})} \sum_{(i,j) \in \Omega_{\text{test}}} r_{ij}$$

$$\bar{r}^{\text{pred}} = \frac{1}{\text{card}(\Omega_{\text{test}})} \sum_{(i,j) \in \Omega_{\text{test}}} r_{ij}^{\text{pred}}$$

```

def SVT(X_input, \
        n_max_iter = 100, \
        threshold = 1000, \
        show=True, \
        tt_split = 0.2, \
        tolerance = 1e-5, \
        seed=seed):

    X_input = X_input.to_numpy()
    mat = X_input.shape

    np.random.seed(seed) # for reproducibility
    indexes = np.random.permutation(mat[0])

    X = X_input.copy()[indexes]

    split = round(tt_split*mat[0])
    X_train = X[split:] # with non-null values
    X_test = X[:split] # with null values

    X_train = X_train.flatten()
    X_test = X_test.flatten()

    train_valid = ~np.isnan(X_train)
    vals_train = X_train[train_valid]

    test_valid = ~np.isnan(X_test)
    vals_test = X_test[test_valid]

    RMSE_SVT_history = []
    pearson_SVT_history = []

    X = np.append(X_train,X_test)
    X[np.isnan(X)] = 0

    X_hat = X.copy()
    for k in range(n_max_iter):
        X_hat_old = X_hat.copy()
        X_hat = X_hat.reshape(mat)
        U,s,VT = np.linalg.svd(X_hat, full_matrices = False)
        s[s < threshold] = 0
        X_hat = U @ np.diag(s) @ VT

        X_hat_train = X_hat[split:].flatten()
        X_hat_test = X_hat[:split].flatten()
        X_hat_train[train_valid] = vals_train

```

```

X_hat = np.append(X_hat_train,X_hat_test)
increment = np.linalg.norm(X_hat - X_hat_old)

vals_test_SVT = X_hat_test[test_valid]
RMSE_SVT = np.sqrt(np.mean((vals_test - vals_test_SVT)**2))
pearson_SVT = pearsonr(vals_test, vals_test_SVT)[0]

RMSE_SVT_history.append(RMSE_SVT)
pearson_SVT_history.append(pearson_SVT)

if show:
    print('==== iteration %d (increment %1.2e)' % (k, increment))
    print('RMSE = %f' % RMSE_SVT)
    print('rho = %f' % pearson_SVT)

if tolerance > increment:
    break

X_hat = X_hat.reshape(mat)
X_hat_train = X_hat[split:].flatten()
X_hat_test = X_hat[:split].flatten()

X_hat_train[train_valid] = vals_train
X_hat_test[test_valid] = vals_test

X_hat = np.append(X_hat_train,X_hat_test)
X_hat = X_hat.reshape(mat)

X_hat = X_hat[-indexes]
X_hat = pd.DataFrame(data=X_hat,columns=avoid_categoricals)
X_hat['age'] = X_hat['age'].astype(int)

if show:
    fig, axs = plt.subplots(2,1,figsize = (8,8))
    axs[0].loglog(RMSE_SVT_history, 'o-')
    axs[1].semilogx(pearson_SVT_history, 'o-')

return X_hat

X_train_cat = X_train[categorical_variables_names]
X_train_avoid = X_train[avoid_categoricals]

```

In order to determine a value for `threshold` is recommended to plot the *singular values* first, to better understand the order of magnitude for the thresholding.


```

X_train_singular = X_train_avoid.to_numpy().copy()
X_train_singular[np.isnan(X_train_avoid)] = 0

_,s,_ = np.linalg.svd(X_train_singular,full_matrices=False)
plt.semilogy(s,'o')

[<matplotlib.lines.Line2D at 0x31c9e59d0>]

```



```

# Train Dataset SVT
X_train_prev = X_train_avoid
X_train_avoid = SVT(X_train_avoid,threshold=2e3,show=False)

print(f"Values for \"age\" previous to SVT: {(X_train_prev['age'] == 65).sum()}")
print(f"Values for \"age\" posterior to SVT: {(X_train_avoid['age'] == 65).sum()}")

Values for "age" previous to SVT: 748
Values for "age" posterior to SVT: 748

for i in X_train_prev.columns.tolist():
    print(f"Variable {i}: \n \
          Previous nulls: {X_train_prev[i].isnull().sum()}\n \
          Posterior nulls: {X_train_avoid[i].isnull().sum()}\n\n")

```



```
Variable age:
  Previous nulls: 0
  Posterior nulls: 0

Variable barthel:
  Previous nulls: 22510
  Posterior nulls: 0

Variable charlson:
  Previous nulls: 59
  Posterior nulls: 0

Variable creatinina:
  Previous nulls: 6322
  Posterior nulls: 0

Variable drg:
  Previous nulls: 2891
  Posterior nulls: 0

Variable estancias:
  Previous nulls: 0
  Posterior nulls: 0

Variable glucosa:
  Previous nulls: 7360
  Posterior nulls: 0

Variable hematocrito:
  Previous nulls: 6588
  Posterior nulls: 0

Variable leucocitos:
  Previous nulls: 6588
  Posterior nulls: 0

Variable num_grupoact3_HOSP:
  Previous nulls: 16768
  Posterior nulls: 0

Variable numurgenciasprevias:
  Previous nulls: 448
```

Posterior nulls: 0

Variable potasio:
Previous nulls: 6872
Posterior nulls: 0

Variable proteina_c_reactiva:
Previous nulls: 14522
Posterior nulls: 0

Variable rdw_cv:
Previous nulls: 6612
Posterior nulls: 0

Variable rdw_sd:
Previous nulls: 6612
Posterior nulls: 0

Variable sodio:
Previous nulls: 6463
Posterior nulls: 0

Variable urea:
Previous nulls: 11344
Posterior nulls: 0

```
X_train =  
pd.DataFrame(data=X_train_cat.to_numpy(),columns=X_train_cat.columns)  
X_train =  
pd.concat([X_train,pd.DataFrame(data=X_train_avoid.to_numpy(),columns=  
X_train_avoid.columns)],axis=1)  
print(X_train.isnull().sum())
```

codidiagingresso	0
codservicioreal	0
metastatic_tumor	0
age	0
barthel	0
charlson	0
creatinina	0
drg	0
estancias	0
glucosa	0

```

hematocrito      0
leucocitos       0
num_grupoact3_HOSP  0
numurgenciasprevias  0
potasio          0
proteina_c_reactiva  0
rdw_cv           0
rdw_sd           0
sodio            0
urea             0
dtype: int64

for c in X_train.columns.tolist():
    median_train_c = X_train[c].apply(pd.to_numeric,
errors='coerce').median()

    X_test[c].fillna(median_train_c, inplace=True)

/Users/ermarinellone/Library/Mobile
Documents/com~apple~CloudDocs/BDS/myenv/lib/python3.11/site-packages/
numpy/lib/nanfunctions.py:1215: RuntimeWarning: Mean of empty slice
    return np.nanmean(a, axis, out=out, keepdims=keepdims)

X_test.isnull().sum()

age              0
barthel          0
charlson         0
codidiagingreso  0
codservicioreal  0
creatinina       0
drg              0
estancias        0
glucosa          0
hematocrito      0
leucocitos       0
metastatic_tumor  0
num_grupoact3_HOSP  0
numurgenciasprevias  0
potasio          0
proteina_c_reactiva  0
rdw_cv           0
rdw_sd           0
sodio            0
urea             0
dtype: int64

```

To fill continuous variables only we opted for an optimal approach, the SVT algorithm, described in the Markdown above.

We used the dataset selecting to avoid the categorical variables, already filled in the previous section, not employed in the matrix given as argument to the function `SVT`.

```
categorical_cols = ['codidiagingreso',  
                    'codservicioreal', 'metastatic_tumor']  
  
combined_data = pd.concat([X_train, X_test], axis=0)  
  
combined_data_encoded = pd.get_dummies(combined_data,  
columns=categorical_cols, prefix=categorical_cols)  
  
X_train_encoded = combined_data_encoded.iloc[:len(X_train),:]  
X_test_encoded = combined_data_encoded.iloc[len(X_train):,:]
```

To manage numerically the string values, we have used an *encoding-like* approach, thanks to the function `pd.get_dummies`.

Merging `X_train` and `X_test` into a single `combined_data` dataset ensures that the one-hot encoding takes into account all categories present in both sets, without having `np.shape` inconsistencies between Train Set and Test Set, otherwise issues could occur during prediction of the values. In this way all the features are represented both in the Train and Test Set, this will be discussed in *sub-section 4*.

NOTE: The previous cell should be in the *sub-section 3.2*, but for simplicity of coding we have left it as the last cell of the pre-processing phase.

4. Categorical Test-Train Dilemma

With our *features extension* approach we have proved that at most one of the new features generated by `pd.get_dummies` have a positive binary value (1), and the rest is set to negative (0).

This is particularly useful when it comes to the Test Set, indeed the unseen features in the Training Set will have the default weight (usually 0, but could be random) applied by the specific *Machine Learning Model* to features constantly null.

Therefore it does not try to improve it from the default weight and it will be kept constantly default until the end of the training!

Hence it is possible to say that the added value will not modify the behaviour of the algorithm, and in the inference phase the weight applied to the unseen feature is nothing but the default weight, since the original model has not been trained on it having a value different from 0.

Block III: Random Forest

```
from sklearn.ensemble import RandomForestClassifier  
  
rf_classifier = RandomForestClassifier()
```

```

print("Default parameters of RandomForestClassifier:")
print(rf_classifier.get_params())

Default parameters of RandomForestClassifier:
{'bootstrap': True, 'ccp_alpha': 0.0, 'class_weight': None,
 'criterion': 'gini', 'max_depth': None, 'max_features': 'sqrt',
 'max_leaf_nodes': None, 'max_samples': None, 'min_impurity_decrease':
 0.0, 'min_samples_leaf': 1, 'min_samples_split': 2,
 'min_weight_fraction_leaf': 0.0, 'n_estimators': 100, 'n_jobs': None,
 'oob_score': False, 'random_state': None, 'verbose': 0, 'warm_start':
 False}

rf_classifier.fit(X_train_encoded, y_train)

RandomForestClassifier()

probabilities = rf_classifier.predict_proba(X_test_encoded)

positive_class_probabilities = probabilities[:, 1]

print(positive_class_probabilities)

[0.25 0.75 0.16 ... 0.24 0.42 0.23]

positive_probs = rf_classifier.predict_proba(X_test_encoded)[:, 1]

auc_roc_rf = round(roc_auc_score(y_test, positive_probs), 4)
print("AUC-ROC Score:", auc_roc_rf)

AUC-ROC Score: 0.7543

```

The AUC-ROC score (Area Under the Receiver Operating Characteristic Curve) is a numerical measure of the area under the ROC curve. This value represents the ability of a classification model to distinguish between positive and negative classes.

In other words, the AUC-ROC score quantifies how well the model is able to separate positive and negative instances without being affected by the specific choice of classification threshold.

The AUC-ROC score is particularly useful when comparing binary classification models and wanting an overall assessment of their performance without having to choose a specific decision threshold.

A higher AUC-ROC value generally indicates a better model in distinguishing between positive and negative classes.

AUC-ROC = 1: Indicates a perfect model that is able to distinguish perfectly between positive and negative classes while AUC-ROC < 0.5: Indicates that the model is performing worse than a random case.

```

predictions = rf_classifier.predict(X_test_encoded)

```

```

accuracy_rf = round(accuracy_score(y_test, predictions),4)
print("Accuracy:", accuracy_rf)

conf_matrix = confusion_matrix(y_test, predictions)

true_negative, false_positive, false_negative, true_positive =
conf_matrix.ravel()

sensitivity = round(true_positive / (true_positive +
false_negative),4)
specificity = round(true_negative / (true_negative +
false_positive),4)

print("Sensitivity (True Positive Rate):", sensitivity)
print("Specificity (True Negative Rate):", specificity)

Accuracy: 0.7185
Sensitivity (True Positive Rate): 0.3835
Specificity (True Negative Rate): 0.8969

```

1) Sensitivity measures the proportion of true positives (correctly classified positive instances) to the total number of true positive instances in the dataset. In other words, it indicates the model's ability to correctly identify instances of the positive class. Sensitivity (TPR) = $TP / (TP + FN)$ Sensitivity is important when the cost of a false negative (a negative prediction for an actually positive instance) is high, for example, in a diagnostic test where missing a positive condition is serious.


2) Specificity measures the proportion of true negatives (correctly classified negative instances) to the total number of true negative instances in the dataset. It indicates the model's ability to correctly identify instances of the negative class. Specificity (TNR) = $TN / (TN + FP)$ Specificity is important when the cost of a false positive (a positive prediction for an actually negative instance) is high, for example, in a test to identify dangerous behavior or safety situations.

```

fpr, tpr, thresholds = roc_curve(y_test, positive_probs)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC Curve (area =
%0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()

```



This code is used to visualize the ROC curve (Receiver Operating Characteristic curve) and calculate the AUC-ROC score for a binary classification model.

First the true positive and false positive rates are calculated for the classes predicted by the model, and then these quantities are used to calculate the Area Under the Curve (AUC) using false positive rates (fpr) and true positive rates (tpr).

Gradient Boosting & SVM (Optional)

Gradient Boosting

```
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import GridSearchCV

gb_clf = GradientBoostingClassifier()
print("Default parameters of Gradient Boosting:")
print(gb_clf.get_params())
```

```
Default parameters of Gradient Boosting:
{'ccp_alpha': 0.0, 'criterion': 'friedman_mse', 'init': None,
 'learning_rate': 0.1, 'loss': 'log_loss', 'max_depth': 3,
 'max_features': None, 'max_leaf_nodes': None, 'min_impurity_decrease':
```

```
0.0, 'min_samples_leaf': 1, 'min_samples_split': 2,  
'min_weight_fraction_leaf': 0.0, 'n_estimators': 100,  
'n_iter_no_change': None, 'random_state': None, 'subsample': 1.0,  
'tol': 0.0001, 'validation_fraction': 0.1, 'verbose': 0, 'warm_start':  
False}
```

```
gb_clf.fit(X_train_encoded, y_train)
```

```
GradientBoostingClassifier()
```

```
positive_probs = gb_clf.predict_proba(X_test_encoded)[: , 1]
```

```
auc_roc_gb = round(roc_auc_score(y_test, positive_probs),4)  
print("AUC-ROC Score:", auc_roc_gb)
```

```
AUC-ROC Score: 0.7506
```

```
predictions = gb_clf.predict(X_test_encoded)
```

```
conf_matrix = confusion_matrix(y_test, predictions)
```

```
true_negative, false_positive, false_negative, true_positive =  
conf_matrix.ravel()
```

```
sensitivity = round(true_positive / (true_positive +  
false_negative),4)  
specificity = round(true_negative / (true_negative +  
false_positive),4)
```

```
print("Sensitivity (True Positive Rate):", sensitivity)  
print("Specificity (True Negative Rate):", specificity)
```

```
Sensitivity (True Positive Rate): 0.261  
Specificity (True Negative Rate): 0.9631
```

```
predictions = gb_clf.predict(X_test_encoded)
```

```
accuracy_gb = round(accuracy_score(y_test, predictions),4)  
print(f'Accuracy of model: {accuracy_gb}')
```

```
class_names = ["Negative", "Positive"]  
print(classification_report(y_test, predictions,  
target_names=class_names))
```


```
Accuracy of model: 0.7192
```

	precision	recall	f1-score	support
Negative	0.71	0.96	0.82	5014
Positive	0.79	0.26	0.39	2670

accuracy			0.72	7684
macro avg	0.75	0.61	0.60	7684
weighted avg	0.74	0.72	0.67	7684

```
probabilities = gb_clf.decision_function(X_test_encoded)
fpr, tpr, thresholds = roc_curve(y_test, probabilities)
auc = roc_auc_score(y_test, probabilities)
```

```
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'AUC = {auc:.2f}')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```



SVM

```
from sklearn.svm import SVC

svm_classifier = SVC(probability=True)

print("Default parameters of SVM Classifier:")
print(svm_classifier.get_params())

Default parameters of SVM Classifier:
{'C': 1.0, 'break_ties': False, 'cache_size': 200, 'class_weight':
None, 'coef0': 0.0, 'decision_function_shape': 'ovr', 'degree': 3,
'gamma': 'scale', 'kernel': 'rbf', 'max_iter': -1, 'probability':
True, 'random_state': None, 'shrinking': True, 'tol': 0.001,
'verbose': False}

cut = round(X_test_encoded.shape[0]*1) # this cut (*<1) is needed to
train the dataset on a smaller Training Set size
svm_classifier.fit(X_train_encoded.iloc[:cut,:], y_train[:cut])

SVC(probability=True)

positive_probs = svm_classifier.predict_proba(X_test_encoded)[: , 1]

auc_roc_svm = round(roc_auc_score(y_test, positive_probs),4)
print("AUC-ROC Score:", auc_roc_svm)

AUC-ROC Score: 0.6816

predictions = svm_classifier.predict(X_test_encoded)

conf_matrix = confusion_matrix(y_test, predictions)

true_negative, false_positive, false_negative, true_positive =
conf_matrix.ravel()

sensitivity = true_positive / (true_positive + false_negative)
specificity = true_negative / (true_negative + false_positive)

print("Sensitivity (True Positive Rate):", sensitivity)
print("Specificity (True Negative Rate):", specificity)

Sensitivity (True Positive Rate): 0.0
Specificity (True Negative Rate): 1.0

predictions = svm_classifier.predict(X_test_encoded)

accuracy_svm = round(accuracy_score(y_test, predictions),4)
print(f'Accuracy of model: {accuracy_svm}')


class_names = ["Negative", "Positive"]
```

```
print(classification_report(y_test, predictions,  
target_names=class_names))
```

Accuracy of model: 0.6525

```
probabilities = svm_classifier.decision_function(X_test_encoded)  
fpr, tpr, thresholds = roc_curve(y_test, probabilities)  
auc = roc_auc_score(y_test, probabilities)
```

```
plt.figure(figsize=(8, 6))  
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'AUC = {auc:.2f}')  
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')  
plt.title('ROC Curve')  
plt.legend(loc='lower right')  
plt.show()
```



Results

```
res = {
    'Models': ['Random Forest', 'Gradient Boosting', 'Support Vector
Machine'],
    'Accuracy': [accuracy_rf, accuracy_gb, accuracy_svm],
    'AUC-ROC-Score': [auc_roc_rf, auc_roc_gb, auc_roc_svm],
}

results = pd.DataFrame(res)
display(HTML(results.to_html(index=False)))

<IPython.core.display.HTML object>
```